

01_setup

December 18, 2025

1 Setup and data cleaning

1.1 Notebook setup

```
[14]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sys
sys.executable
```

```
[14]: '/Users/fanghema/Desktop/aaSTAT_5200/STAT_5200_final_project/env/bin/python'
```

1.2 Data Preprocessing

I then load in the data from replication dataset.

```
[2]: data_galvao_et_al = pd.read_csv('../gmo-files/datareturns.txt', sep="\t", header=None)
data_galvao_et_al
```

```
[2]:      0     1     2     3     4     5     6     7     8     9   ...   43 \
0    -0.67  0.00 -0.31  0.01  0.15  0.012  0.33 -0.45 -0.14 -0.20   ... -1.28
1     0.79 -0.26  0.26 -0.08 -0.19  0.012  0.86  0.65 -0.03  0.29   ...  0.34
2     0.63 -0.17 -0.09  0.19 -0.33  0.012  0.96  0.55  0.83  1.29   ...  1.14
3     0.40  0.08 -0.28  0.07 -0.33  0.012 -0.01  0.09  0.32  0.35   ...  1.38
4    -0.63  0.04 -0.16 -0.31  0.14  0.012 -0.15 -0.28  0.26 -0.37   ...  0.18
...   ...   ...   ...   ...   ...   ...   ...   ...   ...
12960  0.07  0.32 -0.18 -0.24 -0.08  0.000 -1.30 -0.04 -0.04 -0.10   ... -0.34
12961  0.39  0.32 -0.24 -0.20 -0.18  0.000  0.23  0.24 -0.06 -0.03   ...  0.02
12962  0.13  0.22  0.61  0.10  0.17  0.000  0.02 -0.22 -0.21 -0.38   ...  0.53
12963 -0.48  0.06  0.23  0.11  0.25  0.000 -0.12 -0.56 -0.39 -0.82   ... -0.30
12964 -0.93  0.48 -0.45 -0.08 -0.22  0.000 -1.41 -1.51 -1.26 -1.22   ... -1.56

        44     45     46     47     48     49     50     51     52
0    -0.91 -0.86 -0.33 -0.25 -0.23 -1.35 -0.52 -0.24 -0.07
1     1.15  0.29  0.43  0.72  0.59  0.71  0.37  0.27 -0.26
2     0.60  1.13  0.80  0.96  0.35  0.99  0.41  0.49  0.07
```

```

3      -0.26  0.61  0.07 -0.16  0.17  1.11  0.03  0.71  0.83
4     -1.27 -0.71 -0.03 -0.42 -0.39 -0.79 -1.08 -0.03 -0.68
...
12960   0.53  0.01 -0.31 -0.02  0.01 -0.03  0.14 -0.08 -0.20
12961   0.21  0.20  0.39  0.67  0.00  0.01  0.46 -0.12 -0.01
12962   0.29  0.21  0.53  0.59  0.44  0.17  0.42  0.11  0.26
12963  -0.10 -0.27 -0.14 -0.68 -0.15  0.16 -0.35 -0.34 -0.23
12964  -0.46 -0.98 -0.32 -0.21 -1.07 -1.26 -1.33 -0.96 -1.09

```

[12965 rows x 53 columns]

Comparing this to the 49-industry dataset from Kenneth French's data library.

```

[3]: industry_portfolios = pd.read_csv(
    '../data/raw/49_Industry_Portfolios_Daily.csv',
    index_col = 0,
    parse_dates=True
)
ff5_factors = pd.read_csv(
    '../data/raw/F-F_Research_Data_5_Factors_2x3_daily.csv',
    index_col = 0,
    parse_dates=True
)

replication_data = (
    ff5_factors
    .merge(industry_portfolios, left_index=True, right_index=True, how='left')
    .loc["1963-07-01":]
)

replication_data

```

```

[3]:          Mkt-RF    SMB    HML    RMW    CMA     RF    Agric   Food   Soda   Beer   \
1963-07-01   -0.67   0.00  -0.34  -0.01   0.16   0.01    0.33  -0.45  -0.14  -0.20
1963-07-02    0.79  -0.26   0.26  -0.07  -0.20   0.01    0.86   0.65  -0.03   0.29
1963-07-03    0.63  -0.17  -0.09   0.18  -0.34   0.01    0.96   0.55   0.83   1.29
1963-07-05    0.40   0.08  -0.27   0.09  -0.34   0.01   -0.01   0.09   0.32   0.35
1963-07-08   -0.63   0.04  -0.18  -0.29   0.14   0.01   -0.15  -0.28   0.26  -0.37
...
2025-10-27    1.17  -0.81  -1.21  -0.23  -1.19   0.02    0.36   0.59   0.30   1.42
2025-10-28    0.18  -0.34  -0.61   0.63  -1.01   0.02   -0.91  -1.06  -0.01  -1.60
2025-10-29   -0.09  -0.98  -0.81   0.30  -1.37   0.02   -0.46  -3.01  -2.64  -3.19
2025-10-30   -1.10   0.00   0.67  -0.27   0.42   0.02   -2.09  -0.15   0.69   0.82
2025-10-31    0.40  -0.07  -0.24  -1.31   0.18   0.02   -0.29  -0.22  -0.06  -0.69

          ... Boxes  Trans  Whlsl  Rtail  Meals  Banks  Insur  RlEst  Fin   \
1963-07-01 ...  -1.28  -0.91  -0.83  -0.33  -0.25  -0.27  -1.35  -0.51  -0.20
1963-07-02 ...   0.34   1.15   0.33   0.43   0.69   0.47   0.71   0.33   0.43

```

```

1963-07-03 ... 1.14 0.60 1.13 0.80 0.96 -0.01 0.99 0.38 0.80
1963-07-05 ... 1.38 -0.26 0.60 0.07 -0.16 0.14 1.11 0.02 0.64
1963-07-08 ... 0.18 -1.27 -0.71 -0.03 -0.42 -0.45 -0.79 -1.06 -0.02
...
2025-10-27 ... 0.11 1.26 0.31 0.56 0.49 0.52 0.33 -0.64 0.91
2025-10-28 ... -0.67 -0.72 -0.93 0.31 -1.57 -0.17 -1.20 -1.14 -0.54
2025-10-29 ... -4.25 -0.13 -0.52 -0.49 -1.19 -1.04 -1.53 -3.46 -1.34
2025-10-30 ... -0.47 0.39 1.36 -1.78 -1.99 0.65 -1.42 -0.67 -0.55
2025-10-31 ... 0.22 0.85 -0.06 3.93 -1.09 0.32 -0.38 0.13 0.46

      Other
1963-07-01 0.14
1963-07-02 0.13
1963-07-03 0.09
1963-07-05 1.02
1963-07-08 -0.38
...
2025-10-27 -0.72
2025-10-28 -1.40
2025-10-29 -1.51
2025-10-30 0.53
2025-10-31 -0.16

```

[15690 rows x 55 columns]

```

[4]: corr_df = pd.DataFrame(
    index = np.arange(data_galvao_et_al.shape[1]),
    columns = replication_data.columns
)

for i in range(data_galvao_et_al.shape[1]):
    for col in replication_data.columns:
        corr = np.corrcoef(
            data_galvao_et_al.iloc[:, i].values,
            replication_data[col][:data_galvao_et_al.shape[0]].values
        )[0, 1]
        corr_df.loc[i, col] = corr

```

```

[5]: columns_map = {}
for i in corr_df.index:
    argmax = corr_df.loc[i].argmax()
    max = corr_df.loc[i].max()
    print(f"column number {i}: {replication_data.columns[argmax]} ({max})")
    columns_map[i] = replication_data.columns[argmax]

print(f"Number of keys: {len(set(columns_map.keys()))}")
print(f"Number of values: {len(set(columns_map.values()))}")

```

column number 0: Mkt-RF (0.9999631176102544)
column number 1: SMB (0.9983317073995388)
column number 2: HML (0.9707746741309701)
column number 3: RMW (0.9695277483402874)
column number 4: CMA (0.996411334405961)
column number 5: RF (0.9765065560785002)
column number 6: Agric (0.9993517050995383)
column number 7: Food (0.9998997012505297)
column number 8: Soda (0.9999736609358353)
column number 9: Beer (0.9999633984001192)
column number 10: Smoke (0.9996248704879636)
column number 11: Toys (0.9986361761212148)
column number 12: Fun (0.9996860847221328)
column number 13: Books (0.9981487902396878)
column number 14: Hshld (0.9999021300411858)
column number 15: Clths (0.9979290120722208)
column number 16: MedEq (0.9998955127299733)
column number 17: Drugs (0.9998955062570346)
column number 18: Chems (0.9999833203244106)
column number 19: Rubbr (0.9989401195566953)
column number 20: Txtls (0.9978286804583132)
column number 21: BldMt (0.9925498232310794)
column number 22: Cnstr (0.9955760178553862)
column number 23: Steel (0.9990462794382083)
column number 24: FabPr (0.9462112861585128)
column number 25: Mach (0.9897169114843933)
column number 26: ElcEq (0.9981289474892405)
column number 27: Autos (0.9990608501760571)
column number 28: Aero (0.9997757459782651)
column number 29: Ships (0.998256024412168)
column number 30: Guns (0.9997896242645089)
column number 31: Gold (0.9985885553277145)
column number 32: Mines (0.9995911312732187)
column number 33: Coal (0.9986400720334226)
column number 34: Oil (0.9996662148490669)
column number 35: Util (0.9998391848589837)
column number 36: Telcm (0.9989044702521956)
column number 37: PerSv (0.9994444860801095)
column number 38: BusSv (0.9905182743345038)
column number 39: Hardw (0.9961413009522394)
column number 40: Chips (0.9974214473872154)
column number 41: LabEq (0.9971031439150029)
column number 42: Paper (0.9696560686178576)
column number 43: Boxes (0.9978022637534089)
column number 44: Trans (0.9996326924037361)
column number 45: Whls1 (0.9983352779355394)
column number 46: Rtai (0.9999036487527976)
column number 47: Meals (0.9933848285355421)

```

column number 48: Banks (0.9989749220253604)
column number 49: Insur (0.9994260713508203)
column number 50: RlEst (0.9459155034739062)
column number 51: Fin (0.9984710426181009)
column number 52: Other (0.984467019976098)
Number of keys: 53
Number of values: 53

```

We check for the two columns that were not matched to any columns in Galvao et al's dataset. This indeed matches their description.

```
[6]: for col in replication_data.columns:
    if col not in columns_map.values():
        print(col)
```

```
Hlth
Softw
```

We update column names and indices for the original dataframe. We also drop Healthcare and Software from our extended dataframe.

```
[7]: replication_data = replication_data.drop(columns=['Hlth', 'Softw'])
data_galvao_et_al.rename(columns = columns_map, inplace=True)
data_galvao_et_al.index = replication_data.index[:data_galvao_et_al.shape[0]]
```

Data cleaning and replacing invalid values with 0.

```
[8]: replication_data = (
    replication_data
    .replace([np.inf, -np.inf], np.nan)
    .fillna(0)
)

data_galvao_et_al = (
    data_galvao_et_al
    .replace([np.inf, -np.inf], np.nan)
    .fillna(0)
)
```

```
[9]: data_galvao_et_al.to_csv('../data/processed/data_galvao.csv')
replication_data.to_csv("../data/processed/data_extended.csv")
```

1.3 Summary Statistics

```
[11]: replication_data.describe()
```

| | Mkt-RF | SMB | HML | RMW | CMA | \ |
|-------|--------------|--------------|--------------|--------------|--------------|---|
| count | 15690.000000 | 15690.000000 | 15690.000000 | 15690.000000 | 15690.000000 | |
| mean | 0.028841 | 0.005815 | 0.014004 | 0.012736 | 0.011643 | |
| std | 1.023264 | 0.551013 | 0.585571 | 0.403702 | 0.382175 | |

| | | | | | | | |
|-------|--------------|--------------|--------------|--------------|--------------|-----|---|
| min | -17.440000 | -11.150000 | -5.030000 | -2.970000 | -5.310000 | | |
| 25% | -0.420000 | -0.280000 | -0.240000 | -0.180000 | -0.180000 | | |
| 50% | 0.050000 | 0.020000 | 0.010000 | 0.010000 | 0.010000 | | |
| 75% | 0.510000 | 0.310000 | 0.260000 | 0.190000 | 0.200000 | | |
| max | 11.360000 | 6.080000 | 6.730000 | 4.570000 | 2.480000 | | |
| | RF | Agric | Food | Soda | Beer | \ | |
| count | 15690.000000 | 15690.000000 | 15690.000000 | 15690.000000 | 15690.000000 | | |
| mean | 0.017183 | 0.049493 | 0.044612 | 0.053917 | 0.049760 | | |
| std | 0.012740 | 1.437237 | 0.905245 | 1.356966 | 1.126727 | | |
| min | 0.000000 | -15.270000 | -16.080000 | -19.220000 | -14.730000 | | |
| 25% | 0.010000 | -0.630000 | -0.410000 | -0.610000 | -0.530000 | | |
| 50% | 0.020000 | 0.040000 | 0.060000 | 0.050000 | 0.040000 | | |
| 75% | 0.020000 | 0.730000 | 0.510000 | 0.690000 | 0.620000 | | |
| max | 0.060000 | 20.320000 | 10.010000 | 11.680000 | 11.460000 | | |
| | ... | Boxes | Trans | Whlsl | Rtail | \ | |
| count | ... | 15690.000000 | 15690.000000 | 15690.000000 | 15690.000000 | | |
| mean | ... | 0.045228 | 0.045016 | 0.048480 | 0.052345 | | |
| std | ... | 1.260816 | 1.252522 | 1.088046 | 1.156068 | | |
| min | ... | -21.350000 | -17.570000 | -13.250000 | -17.970000 | | |
| 25% | ... | -0.610000 | -0.600000 | -0.480000 | -0.520000 | | |
| 50% | ... | 0.050000 | 0.050000 | 0.070000 | 0.070000 | | |
| 75% | ... | 0.700000 | 0.680000 | 0.610000 | 0.620000 | | |
| max | ... | 10.910000 | 12.710000 | 10.640000 | 11.750000 | | |
| | ... | Meals | Banks | Insur | RlEst | Fin | \ |
| count | 15690.000000 | 15690.000000 | 15690.000000 | 15690.000000 | 15690.000000 | | |
| mean | 0.054479 | 0.047213 | 0.047164 | 0.035504 | 0.055850 | | |
| std | 1.247532 | 1.417553 | 1.174384 | 1.449594 | 1.427951 | | |
| min | -15.260000 | -16.960000 | -15.200000 | -17.640000 | -16.360000 | | |
| 25% | -0.577500 | -0.540000 | -0.490000 | -0.600000 | -0.500000 | | |
| 50% | 0.070000 | 0.040000 | 0.060000 | 0.050000 | 0.070000 | | |
| 75% | 0.690000 | 0.630000 | 0.600000 | 0.680000 | 0.640000 | | |
| max | 15.890000 | 16.940000 | 17.820000 | 17.120000 | 18.180000 | | |
| | ... | Other | | | | | |
| count | 15690.000000 | | | | | | |
| mean | 0.030916 | | | | | | |
| std | 1.408428 | | | | | | |
| min | -17.230000 | | | | | | |
| 25% | -0.590000 | | | | | | |
| 50% | 0.050000 | | | | | | |
| 75% | 0.680000 | | | | | | |
| max | 16.840000 | | | | | | |

[8 rows x 53 columns]

```
[12]: data_galvao_et_al.describe()
```

| | Mkt-RF | SMB | HML | RMW | CMA | \ |
|-------|------------------|--------------|--------------|--------------|--------------|---|
| count | 12965.000000 | 12965.000000 | 12965.000000 | 12965.000000 | 12965.000000 | |
| mean | 0.024251 | 0.009328 | 0.017895 | 0.012640 | 0.015580 | |
| std | 0.992087 | 0.520785 | 0.495309 | 0.348898 | 0.360934 | |
| min | -17.440000 | -11.180000 | -5.090000 | -2.860000 | -5.240000 | |
| 25% | -0.420000 | -0.260000 | -0.210000 | -0.160000 | -0.170000 | |
| 50% | 0.050000 | 0.030000 | 0.010000 | 0.010000 | 0.010000 | |
| 75% | 0.490000 | 0.300000 | 0.240000 | 0.170000 | 0.190000 | |
| max | 11.350000 | 6.090000 | 4.040000 | 4.100000 | 2.490000 | |
| | | | | | | |
| | RF | Agric | Food | Soda | Beer | \ |
| count | 12965.000000 | 12965.000000 | 12965.000000 | 12965.000000 | 12965.000000 | |
| mean | 0.019350 | 0.051073 | 0.050685 | 0.056383 | 0.054464 | |
| std | 0.012195 | 1.396318 | 0.879694 | 1.406782 | 1.128794 | |
| min | 0.000000 | -15.270000 | -16.040000 | -19.220000 | -14.730000 | |
| 25% | 0.012000 | -0.630000 | -0.390000 | -0.630000 | -0.530000 | |
| 50% | 0.019000 | 0.040000 | 0.060000 | 0.050000 | 0.040000 | |
| 75% | 0.025000 | 0.730000 | 0.500000 | 0.720000 | 0.620000 | |
| max | 0.061000 | 20.320000 | 9.980000 | 11.680000 | 10.120000 | |
| | | | | | | |
| | Boxes | Trans | Whlsl | Rtail | \ | |
| count | ... 12965.000000 | 12965.000000 | 12965.000000 | 12965.000000 | | |
| mean | ... 0.048595 | 0.046272 | 0.048008 | 0.050200 | | |
| std | ... 1.243543 | 1.215566 | 1.059366 | 1.133851 | | |
| min | ... -21.430000 | -17.560000 | -13.200000 | -18.000000 | | |
| 25% | ... -0.600000 | -0.590000 | -0.470000 | -0.520000 | | |
| 50% | ... 0.050000 | 0.050000 | 0.080000 | 0.060000 | | |
| 75% | ... 0.680000 | 0.670000 | 0.590000 | 0.600000 | | |
| max | ... 10.920000 | 9.330000 | 9.750000 | 11.750000 | | |
| | | | | | | |
| | Meals | Banks | Insur | RlEst | Fin | \ |
| count | 12965.000000 | 12965.000000 | 12965.000000 | 12965.000000 | 12965.000000 | |
| mean | 0.055287 | 0.044963 | 0.046565 | 0.029755 | 0.054201 | |
| std | 1.252471 | 1.382958 | 1.153460 | 1.431005 | 1.414357 | |
| min | -15.480000 | -16.980000 | -13.980000 | -15.500000 | -16.280000 | |
| 25% | -0.590000 | -0.520000 | -0.490000 | -0.600000 | -0.480000 | |
| 50% | 0.070000 | 0.040000 | 0.060000 | 0.030000 | 0.060000 | |
| 75% | 0.720000 | 0.590000 | 0.580000 | 0.680000 | 0.610000 | |
| max | 11.490000 | 16.960000 | 17.840000 | 21.900000 | 17.940000 | |
| | | | | | | |
| | Other | | | | | |
| count | 12965.000000 | | | | | |
| mean | 0.029523 | | | | | |
| std | 1.433517 | | | | | |
| min | -17.260000 | | | | | |

```

25%      -0.610000
50%      0.050000
75%      0.700000
max      16.840000

```

[8 rows x 53 columns]

```

[15]: factors = ["Mkt-RF", "SMB", "HML", "RMW", "CMA"]

plt.figure(figsize=(14, 10))

for i, fac in enumerate(factors, 1):
    plt.subplot(3, 2, i)
    replication_data[fac].plot(kind="kde", label="Extended Sample", linewidth=2)
    data_galvao_et_al[fac].plot(kind="kde", label="Original Sample", linewidth=2)
    plt.title(f"{fac} - KDE Comparison")
    plt.xlabel("Daily Return")
    plt.legend()

plt.tight_layout()
plt.show()

combined = pd.concat([
    replication_data[factors].assign(sample="Extended"),
    data_galvao_et_al[factors].assign(sample="Original")
])

combined_melted = combined.melt(id_vars="sample")

plt.figure(figsize=(12, 6))
sns.boxplot(data=combined_melted, x="variable", y="value", hue="sample")
plt.title("Factor Return Distribution Comparison (Boxplots)")
plt.xlabel("Factor")
plt.ylabel("Daily Return")
plt.show()

plt.figure(figsize=(14, 10))

for i, fac in enumerate(factors, 1):
    plt.subplot(3, 2, i)
    replication_data[fac].rolling(252).std().plot(label="Extended Sample", linewidth=1.5)
    data_galvao_et_al[fac].rolling(252).std().plot(label="Original Sample", linewidth=1.5)
    plt.title(f"{fac} - Rolling 1-Year Volatility")
    plt.xlabel("Date")

```

```

plt.ylabel("Volatility")
plt.legend()

plt.tight_layout()
plt.show()


stats = pd.DataFrame({
    "mean_diff": replication_data[factors].mean() - data_galvao_et_al[factors].
    mean(),
    "std_diff": replication_data[factors].std() - data_galvao_et_al[factors].
    std(),
    "skew_diff": replication_data[factors].skew() - data_galvao_et_al[factors].
    skew(),
    "kurt_diff": replication_data[factors].kurt() - data_galvao_et_al[factors].
    kurt()
})



```


```

plt.figure(figsize=(8, 5))
sns.heatmap(stats, annot=True, cmap="coolwarm", center=0, fmt=".4f")
plt.title("Differences in Statistical Moments (Extended - Original)")
plt.ylabel("Factor")
plt.show()

```







```
[16]: N_SAMPLE_INDUSTRIES = 6

factor_cols = ["Mkt-RF", "SMB", "HML", "RMW", "CMA", "RF"]
industry_cols = [col for col in replication_data.columns if col not in factor_cols]

np.random.seed(42)
sampled_industries = np.random.choice(industry_cols, size=N_SAMPLE_INDUSTRIES, replace=False).tolist()

print("Sampled industries:", sampled_industries)

plt.figure(figsize=(14, 10))

for i, ind in enumerate(sampled_industries, 1):
    plt.subplot(3, 3, i)
    replication_data[ind].plot(kind="kde", label="Extended Sample", linewidth=2)
    data_galvao_et_al[ind].plot(kind="kde", label="Original Sample", linewidth=2)
    plt.title(f"{ind} - KDE Comparison")
    plt.xlabel("Daily Industry Return")
    plt.legend()
```

```

plt.tight_layout()
plt.show()

combined_ind = pd.concat([
    replication_data[sampled_industries].assign(sample="Extended"),
    data_galvao_et_al[sampled_industries].assign(sample="Original")
])

combined_ind_melted = combined_ind.melt(id_vars="sample")

plt.figure(figsize=(14, 6))
sns.boxplot(data=combined_ind_melted, x="variable", y="value", hue="sample")
plt.title("Industry Return Distribution Comparison (Boxplots)")
plt.xlabel("Industry Portfolio")
plt.ylabel("Daily Return")
plt.xticks(rotation=45)
plt.show()

plt.figure(figsize=(14, 10))

for i, ind in enumerate(sampled_industries, 1):
    plt.subplot(3, 3, i)
    replication_data[ind].rolling(252).std().plot(label="Extended Sample", u
    ↵linewidth=1.5)
    data_galvao_et_al[ind].rolling(252).std().plot(label="Original Sample", u
    ↵linewidth=1.5)
    plt.title(f"{ind} - Rolling 1Y Volatility")
    plt.xlabel("Date")
    plt.ylabel("Volatility")
    plt.legend()

plt.tight_layout()
plt.show()

stats_ind = pd.DataFrame({
    "mean_diff": replication_data[sampled_industries].mean() - u
    ↵data_galvao_et_al[sampled_industries].mean(),
    "std_diff": replication_data[sampled_industries].std() - u
    ↵data_galvao_et_al[sampled_industries].std(),
    "skew_diff": replication_data[sampled_industries].skew() - u
    ↵data_galvao_et_al[sampled_industries].skew(),
    "kurt_diff": replication_data[sampled_industries].kurt() - u
    ↵data_galvao_et_al[sampled_industries].kurt()
})

```

)

```
plt.figure(figsize=(10, 6))
sns.heatmap(stats_ind, annot=True, cmap="coolwarm", center=0, fmt=".4f")
plt.title("Differences in Industry Return Moments (Extended - Original)")
plt.ylabel("Industry")
plt.show()
```

Sampled industries: ['Coal', 'Whlsl', 'Mines', 'Insur', 'Guns', 'Paper']

